

Robust and efficient methods for single-cell CRISPR screens: PerturbPlan and spaCRT

McGill Statistics Seminar
April 2026

Ziang Niu
University of Pennsylvania
<https://ziangniu6.github.io>


A model for genetic origins of human diseases

A model for genetic origins of human diseases

Normal


A model for genetic origins of human diseases

Normal



A model for genetic origins of human diseases

Normal




A model for genetic origins of human diseases


Normal




A model for genetic origins of human diseases




A model for genetic origins of human diseases




A model for genetic origins of human diseases




A model for genetic origins of human diseases





A model for genetic origins of human diseases





A model for genetic origins of human diseases





Uncovering genetic origins of human diseases



Uncovering genetic origins of human diseases





Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases




Single-cell CRISPR screens

Uncovering genetic origins of human diseases




Single-cell CRISPR screens

Uncovering genetic origins of human diseases





CRISPR





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases





Single-cell CRISPR screens

Uncovering genetic origins of human diseases



Single-cell CRISPR screens


Uncovering genetic origins of human diseases




Single-cell CRISPR screens

Clinical promise for monogenic disease

Clinical promise for monogenic disease




Clinical promise for monogenic disease




- CASGEVY: the first-ever FDA-approved gene editing therapy.

Clinical promise for monogenic disease



- CASGEVY: the first-ever FDA-approved gene editing therapy.
- Apply CRISPR base-editing to edit a key erythroid enhancer (of BCL11A gene) in sickle cell disease to restore blood hemoglobin to normal levels.

Clinical promise for monogenic disease




- CASGEVY: the first-ever FDA-approved gene editing therapy.
- Apply CRISPR base-editing to edit a key erythroid enhancer (of BCL11A gene)¹ in sickle cell disease to restore blood hemoglobin to normal levels.

¹Bauer et al (2013)

Perturb-seq protocol

Perturb-seq protocol

Infect cells with CRISPR gRNAs
targeting genomic elements




Single-cell
RNA sequencing



For each cell, measure:
1. gRNA expression
2. Gene expression

Perturb-seq protocol

Infect cells with CRISPR gRNAs targeting genomic elements




Single-cell RNA sequencing



For each cell, measure:
1. gRNA expression
2. Gene expression

Test element-gene pairs for regulatory relationships.



Bottlenecks: scale, cost and sparsity

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹

¹Replogle et al (2022)

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,

¹Replogle et al (2022)

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.

¹Replogle et al (2022)

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.



Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.

\$180k

Experimental challenge: high expense and under-power without careful design.

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.

\$180k

Experimental challenge: high expense and under-power without careful design.

Analytical challenge: computational burden.

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.

\$180k

Experimental challenge: high expense and under-power without careful design.

Analytical challenge: computational burden.

Cell perturbed	Gene Expression					
	0	1	2	3	4	8
FALSE	201027	5594	324	32	5	1
TRUE	332	9	0	0	0	0

Bottlenecks: scale, cost and sparsity

Genome-wide gene-targeting:¹
2M cells, 10k perturbations,
10k genes, 11k test pairs.

\$180k

Experimental challenge: high expense and under-power without careful design.

Analytical challenge: computational burden.

Cell perturbed	Gene Expression					
	0	1	2	3	4	8
FALSE	201027	5594	324	32	5	1
TRUE	332	9	0	0	0	0

Analytical challenge: classical CLT-based test (Wald/Score) fails to control Type-I error.

Bottlenecks: scale, cost and sparsity

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Solutions: PerturbPlan and spaCRT

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Solutions: PerturbPlan and spaCRT

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Prospective experimental design: **PerturbPlan** balances cost and power in Perturb-seq experiment, supporting 11 common design problems.

Solutions: PerturbPlan and spaCRT

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Prospective experimental design: **PerturbPlan** balances cost and power in Perturb-seq experiment, supporting 11 common design problems.




Solutions: PerturbPlan and spaCRT

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Prospective experimental design: **PerturbPlan** balances cost and power in Perturb-seq experiment, supporting 11 common design problems.




Solutions: PerturbPlan and spaCRT

Experimental challenge: high expense + under-power without careful design.

Analytical challenge: computational burden + failure of classical asymptotic test.

Prospective experimental design: **PerturbPlan** balances cost and power in Perturb-seq experiment, supporting 11 common design problems.


Post-experimental analysis: **spaCRT** supports computationally efficient differential expression analysis, with strong statistical validity guarantee.




Part I: PerturbPlan

Designing a Perturb-seq experiment is hard

Designing a Perturb-seq experiment is hard




Designing a Perturb-seq experiment is hard




Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4
.....			

Designing a Perturb-seq experiment is hard




Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4
.....			

Designing a Perturb-seq experiment is hard




Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4
.....			

Designing a Perturb-seq experiment is hard



Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4
.....			

Designing a Perturb-seq experiment is hard



Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4


.....

Design objectives

Cost: Keep library prep and sequencing costs under budget.

Power: Discover large fraction of true regulatory relationships.

Designing a Perturb-seq experiment is hard



Cost (\$)	#cells/perturb	#reads/cell	Power
10k	1k	10k	0.7
10k	500	20k	0.6
10k	250	50k	0.4

.....

Design objectives

Cost: Keep library prep and sequencing costs under budget.

Power: Discover large fraction of true regulatory relationships.

Technical challenge

Compute power quickly and accurately.

A satisfactory solution is not yet available

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]

- ▶ Simulations are computationally expensive

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]
 - Simulations are computationally expensive
2. Sequential identification of perturb-seq targets
[Huang et al 2023, Roohani 2025, Hao et al 2025]

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]

- Simulations are computationally expensive

2. Sequential identification of perturb-seq targets
[Huang et al 2023, Roohani 2025, Hao et al 2025]

- Does not address power or cost

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]

- Simulations are computationally expensive

2. Sequential identification of perturb-seq targets
[Huang et al 2023, Roohani 2025, Hao et al 2025]

- Does not address power or cost

3. Power analysis for single-cell differential expression
[Vieth 2017, Crowell 2020, Schmid 2021]

A satisfactory solution is not yet available

1. Simulation-based analysis of power for perturb-seq
[Ray 2025]

- ▶ Simulations are computationally expensive

2. Sequential identification of perturb-seq targets
[Huang et al 2023, Roohani 2025, Hao et al 2025]

- ▶ Does not address power or cost

3. Power analysis for single-cell differential expression
[Vieth 2017, Crowell 2020, Schmid 2021]

- ▶ Not directly applicable to Perturb-seq

We propose an analytical power formula

We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{ hypotheses with signal}$

We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$

$\approx$ power of detecting element l on gene j

We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$

$\approx \mathbb{E}_{(j,l) \sim E}[\text{power of detecting element } l \text{ on gene } j]$


Element-gene pairs

The diagram illustrates the relationship between the expectation formula and the element-gene pairs. A rounded rectangle labeled "Element-gene pairs" is connected by an arrow to the term (j,l) in the expectation formula $\mathbb{E}_{(j,l) \sim E}[\text{power of detecting element } l \text{ on gene } j]$. The arrow points from the rectangle to the (j,l) term, indicating that the expectation is taken over the distribution of element-gene pairs.

We propose an analytical power formula


Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$

Element-
gene pairs

$$\approx \mathbb{E}_{(j,l) \sim E}[\text{power of detecting element } l \text{ on gene } j]$$
$$\approx \mathbb{E}_{(j,l) \sim E}[\Phi(z_{\hat{t}} - h_{jl})]$$



We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$




We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$




We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$




We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$




We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$



We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$




(α_j, θ_j) : relative expression and dispersion of gene j

$(\beta_{jl}, \text{sd}_g)$: effect size for element l on gene j and gRNA variability


We propose an analytical power formula

Overall power = $\mathbb{E}[\# \text{ true discoveries}] / \# \text{hypotheses with signal}$





Formula matches simulations


Formula matches simulations



Formula matches simulations but is much faster



Formula matches simulations but is much faster




PerturbPlan tracks downsampling on real data

PerturbPlan tracks downsampling on real data




Perturb-seq
data



PerturbPlan tracks downsampling on real data





PerturbPlan tracks downsampling on real data




PerturbPlan tracks downsampling on real data



PerturbPlan tracks downsampling on real data




Positive control fold change




Downsampling PerturbPlan (max FC) PerturbPlan (median FC)

PerturbPlan web app facilitates interactive design


PerturbPlan web app facilitates interactive design




PerturbPlan web app facilitates interactive design




PerturbPlan web app facilitates interactive design




PerturbPlan web app facilitates interactive design




PerturbPlan web app facilitates interactive design



PerturbPlan web app facilitates interactive design



PerturbPlan web app facilitates interactive design



PerturbPlan addresses 11 design problems

Minimize

- Cells per target
- Reads per cell
- Expression threshold
- Fold change
- ✓ Total cost

Example: Find the minimum cost for which power is at least 80%, varying cells per target and reads per cell.

Subject to constraints on

- ✓ Power
- Cost

While varying

- ✓ Cells per target
- ✓ Reads per cell

PerturbPlan addresses 11 design problems

Minimize

- Cells per target
- Reads per cell
- Expression threshold
- Fold change
- ✓ Total cost


Subject to constraints on

- ✓ Power
- Cost

While varying


- ✓ Cells per target
- ✓ Reads per cell

Example: Find the minimum cost for which power is at least 80%, varying cells per target and reads per cell.




Part II: spaCRT

Perturbation-gene association testing



Perturbation-gene association testing


Focus on one genomic element and one gene.



Perturbation-gene association testing

Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,




Perturbation-gene association testing

Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,

- $X_i \in \{0,1\}$: perturbation presence




Perturbation-gene association testing

Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,

- $X_i \in \{0,1\}$: perturbation presence
- $Y_i \in \mathbb{N}$: gene expression




Perturbation-gene association testing

Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,

- $X_i \in \{0,1\}$: perturbation presence
- $Y_i \in \mathbb{N}$: gene expression
- $Z_i \in \mathbb{R}^d$: Vector of covariates




Perturbation-gene association testing

Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,

- $X_i \in \{0,1\}$: perturbation presence
- $Y_i \in \mathbb{N}$: gene expression
- $Z_i \in \mathbb{R}^d$: Vector of covariates

The goal is to test



Perturbation-gene association testing


Focus on one genomic element and one gene.

For each cell $i = 1, \dots, n$,

- $X_i \in \{0, 1\}$: perturbation presence
- $Y_i \in \mathbb{N}$: gene expression
- $Z_i \in \mathbb{R}^d$: Vector of covariates

The goal is to test

$$H_0 : X_i \perp\!\!\!\perp Y_i \mid Z_i.$$



Model for the count-based CRISPR data

Model for the count-based CRISPR data

Working model:
$$\begin{cases} X_i \mid Z_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi_i); & \text{logit}(\pi_i) = Z_i^T \delta \\ Y_i \mid X_i, Z_i \stackrel{\text{ind}}{\sim} \text{NegBin}(\mu_i, \theta); & \log \mu_i = X_i \beta + Z_i^T \gamma \end{cases}.$$

Model for the count-based CRISPR data

Working model:
$$\begin{cases} X_i | Z_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi_i); & \text{logit}(\pi_i) = Z_i^T \delta \\ Y_i | X_i, Z_i \stackrel{\text{ind}}{\sim} \text{NegBin}(\mu_i, \theta); & \log \mu_i = X_i \beta + Z_i^T \gamma \end{cases}.$$

Three tests:

1. Negative binomial regression score test.

Model for the count-based CRISPR data

Working model:
$$\begin{cases} X_i | Z_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi_i); & \text{logit}(\pi_i) = Z_i^T \delta \\ Y_i | X_i, Z_i \stackrel{\text{ind}}{\sim} \text{NegBin}(\mu_i, \theta); & \log \mu_i = X_i \beta + Z_i^T \gamma \end{cases}.$$

Three tests:

1. Negative binomial regression score test.
2. Generalized covariance measure (GCM)¹ test, based on normal limit of
$$T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mathbb{E}}[X_i | Z_i])(Y_i - \hat{\mathbb{E}}[Y_i | Z_i]).$$

¹Shah and Peters (2020)

Model for the count-based CRISPR data

Working model:
$$\begin{cases} X_i | Z_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi_i); & \text{logit}(\pi_i) = Z_i^T \delta \\ Y_i | X_i, Z_i \stackrel{\text{ind}}{\sim} \text{NegBin}(\mu_i, \theta); & \log \mu_i = X_i \beta + Z_i^T \gamma \end{cases}.$$

Three tests:


1. Negative binomial regression score test.
2. Generalized covariance measure (GCM)¹ test, based on normal limit of $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mathbb{E}}[X_i | Z_i])(Y_i - \hat{\mathbb{E}}[Y_i | Z_i])$.
3. Distilled conditional randomization test (dCRT)², based on T_n and resampling from fitted distribution $X_i | Z_i$.

¹Shah and Peters (2020)


²Liu et al. (2022), Barry et al. (2021)

Challenges: scale and sparsity

Challenges: scale and sparsity




Challenges: scale and sparsity



Cell perturbed	Gene Expression					
	0	1	2	3	4	8
FALSE	201027	5594	324	32	5	1
TRUE	332	9	0	0	0	0


Challenges: scale and sparsity



Cell perturbed	Gene Expression					
	0	1	2	3	4	8
FALSE	201027	5594	324	32	5	1
TRUE	332	9	0	0	0	0

Statistical challenge: CLT does not kick in.

Challenges: scale and sparsity




In Gasperini (2019) data¹, we have $n \approx 200,000$ cells, $p \approx 13,000$ genes, and $m \approx 6,000$ genomic elements.

		Gene Expression					
Cell perturbed		0	1	2	3	4	8
FALSE		201027	5594	324	32	5	1
TRUE		332	9	0	0	0	0

Statistical challenge: CLT does not kick in.

¹Gasperini et al (2019)

Challenges: scale and sparsity



In Gasperini (2019) data¹, we have $n \approx 200,000$ cells, $p \approx 13,000$ genes, and $m \approx 6,000$ genomic elements.

Cell perturbed	Gene Expression					
	0	1	2	3	4	8
FALSE	201027	5594	324	32	5	1
TRUE	332	9	0	0	0	0

Statistical challenge: CLT does not kick in.
Computational challenge: directly applying computationally heavy is infeasible.

Challenges: scale and sparsity

Challenges: scale and sparsity

Goal

Efficient computation
under error-control
constraints

Challenges: scale and sparsity

Effective Sample size
(# cells w. nonzero
transcripts)

Multiplicity
(# E-G pairs of
interest)

Goal

Efficient computation
under error-control
constraints

Computation efficiency
(# resamples required)

Challenges: scale and sparsity

Effective Sample size
(# cells w. nonzero
transcripts)

Multiplicity
(# E-G pairs of
interest)


Computation burden
scales with ESS

Goal


Efficient computation
under error-control
constraints

Computation efficiency
(# resamples required)


Challenges: scale and sparsity



Challenges: scale and sparsity




Challenges: scale and sparsity




Benchmarking Type-I error on negative controls


Benchmarking Type-I error on negative controls





Benchmarking Type-I error on negative controls





Benchmarking Type-I error on negative controls





Benchmarking Type-I error on negative controls





Benchmarking Type-I error on negative controls



Benchmarking Type-I error on negative controls



Benchmarking Type-I error on negative controls



The distilled conditional randomization test

dCRT algorithm

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
2. for $b = 1, \dots, B$:

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
2. for $b = 1, \dots, B$:
 - a. Sample $\tilde{X}_i^{(b)} \stackrel{\text{ind}}{\sim} \text{Ber}(\hat{\pi}(Z_i))$.

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
2. for $b = 1, \dots, B$:
 - a. Sample $\tilde{X}_i^{(b)} \stackrel{\text{ind}}{\sim} \text{Ber}(\hat{\pi}(Z_i))$.
 - b. Compute $\tilde{T}_n^{(b)} \equiv \frac{1}{n} \sum_{i=1}^n (\tilde{X}_i^{(b)} - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.

The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
2. for $b = 1, \dots, B$:
 - a. Sample $\tilde{X}_i^{(b)} \stackrel{\text{ind}}{\sim} \text{Ber}(\hat{\pi}(Z_i))$.
 - b. Compute $\tilde{T}_n^{(b)} \equiv \frac{1}{n} \sum_{i=1}^n (\tilde{X}_i^{(b)} - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
3. dCRT p-value $p_n^{\text{dCRT}} = \frac{1}{B+1} \left(1 + \sum_{b=1}^B 1(\tilde{T}_n^{(b)} \geq T_n) \right)$.


The distilled conditional randomization test

Suppose $X_i \stackrel{\text{ind}}{\sim} \text{Ber}(\pi(Z_i))$ for some function π .

Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .


dCRT algorithm

1. Compute $T_n \equiv \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
2. for $b = 1, \dots, B$:
 - a. Sample $\tilde{X}_i^{(b)} \stackrel{\text{ind}}{\sim} \text{Ber}(\hat{\pi}(Z_i))$.
 - b. Compute $\tilde{T}_n^{(b)} \equiv \frac{1}{n} \sum_{i=1}^n (\tilde{X}_i^{(b)} - \hat{\pi}(Z_i))(Y_i - \hat{\mu}(Z_i))$.
3. dCRT p-value $p_n^{\text{dCRT}} = \frac{1}{B+1} \left(1 + \sum_{b=1}^B 1(\tilde{T}_n^{(b)} \geq T_n) \right)$.




Approximating dCRT p -values

Approximating dCRT p -values




Approximating dCRT p -values

- Draw moderate number of resamples and fit parametric curve to distribution.



Approximating dCRT p -values


- Draw moderate number of resamples and fit parametric curve to distribution.¹



¹Barry et al. (2021)

Approximating dCRT p -values


- Draw moderate number of resamples and fit parametric curve to distribution.¹
- Draw an adaptive number of resamples based on test statistic.



¹Barry et al. (2021)

Approximating dCRT p -values

- Draw moderate number of resamples and fit parametric curve to distribution.¹
- Draw an adaptive number of resamples based on test statistic.²




¹Barry et al. (2021)

²Besag and Clifford (1991), Fay et al. (2007), Gandy (2009), Silva and Assuncao (2013), Fischer and Ramdas (2024)

Approximating dCRT p -values

- Draw moderate number of resamples and fit parametric curve to distribution.¹
- Draw an adaptive number of resamples based on test statistic.²
- Use the saddlepoint approximation for sums of independent random variables.




¹Barry et al. (2021)

²Besag and Clifford (1991), Fay et al. (2007), Gandy (2009), Silva and Assuncao (2013), Fischer and Ramdas (2024)

Approximating dCRT p -values

- Draw moderate number of resamples and fit parametric curve to distribution.¹
- Draw an adaptive number of resamples based on test statistic.²
- Use the saddlepoint approximation for sums of independent random variables.³




¹Barry et al. (2021)

²Besag and Clifford (1991), Fay et al. (2007), Gandy (2009), Silva and Assuncao (2013), Fischer and Ramdas (2024)

³Lugannani and Rice (1980), Robinson (1982), Davison and Hinkley (1988), Jing et al. (1994)

Approximating dCRT p -values

- Draw moderate number of resamples and fit parametric curve to distribution.¹
- Draw an adaptive number of resamples based on test statistic.²
- Use the saddlepoint approximation for sums of independent random variables.³



¹Barry et al. (2021)

²Besag and Clifford (1991), Fay et al. (2007), Gandy (2009), Silva and Assuncao (2013), Fischer and Ramdas (2024)

³Lugannani and Rice (1980), Robinson (1982), Davison and Hinkley (1988), Jing et al. (1994)

spaCRT: a resampling-free approximation for dCRT

spaCRT: a resampling-free approximation for dCRT

spaCRT algorithm (Niu et al., 2024)

spaCRT: a resampling-free approximation for dCRT

spaCRT algorithm (Niu et al., 2024)

1. Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

spaCRT: a resampling-free approximation for dCRT

spaCRT algorithm (Niu et al., 2024)

1. Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .
2. Let $\hat{s}_n$ be the solution of the saddlepoint equation in s :

$$\sum_{i=1}^n \left\{ X_i - \text{expit} \left(\text{logit}(\hat{\pi}(Z_i)) + s(Y_i - \hat{\mu}(Z_i)) \right) \right\} \{ Y_i - \hat{\mu}(Z_i) \} = 0.$$

spaCRT: a resampling-free approximation for dCRT

spaCRT algorithm (Niu et al., 2024)

1. Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .
2. Let $\hat{s}_n$ be the solution of the saddlepoint equation in s :

$$\sum_{i=1}^n \left\{ X_i - \text{expit} \left(\text{logit}(\hat{\pi}(Z_i)) + s(Y_i - \hat{\mu}(Z_i)) \right) \right\} \{ Y_i - \hat{\mu}(Z_i) \} = 0.$$

3. Letting $\tilde{\pi}_i = \text{expit} \left(\text{logit}(\hat{\pi}(Z_i)) + \hat{s}_n(Y_i - \hat{\mu}(Z_i)) \right)$ and $\hat{\pi}_i = \hat{\pi}(Z_i)$, define

$$\lambda_n = \hat{s}_n \sqrt{\sum_{i=1}^n (Y_i - \hat{\mu}(Z_i))^2 \tilde{\pi}(1 - \tilde{\pi})}; \quad r_n = \text{sgn}(\hat{s}_n) \sqrt{2 \sum_{i=1}^n \left\{ X_i \log \frac{\tilde{\pi}}{\hat{\pi}} + (1 - X_i) \log \frac{1 - \tilde{\pi}}{1 - \hat{\pi}} \right\}}.$$

spaCRT: a resampling-free approximation for dCRT

spaCRT algorithm (Niu et al., 2024)

1. Learn $\hat{\pi}(Z_i)$ by regressing X on Z and $\hat{\mu}(Z_i) \approx \mathbb{E}[Y_i | Z_i]$ by regressing Y on Z .

2. Let $\hat{s}_n$ be the solution of the saddlepoint equation in s :

$$\sum_{i=1}^n \left\{ X_i - \text{expit} \left(\text{logit}(\hat{\pi}(Z_i)) + s(Y_i - \hat{\mu}(Z_i)) \right) \right\} \{ Y_i - \hat{\mu}(Z_i) \} = 0.$$


3. Letting $\tilde{\pi}_i = \text{expit} \left(\text{logit}(\hat{\pi}(Z_i)) + \hat{s}_n(Y_i - \hat{\mu}(Z_i)) \right)$ and $\hat{\pi}_i = \hat{\pi}(Z_i)$, define

$$\lambda_n = \hat{s}_n \sqrt{\sum_{i=1}^n (Y_i - \hat{\mu}(Z_i))^2 \tilde{\pi}(1 - \tilde{\pi})}; \quad r_n = \text{sgn}(\hat{s}_n) \sqrt{2 \sum_{i=1}^n \left\{ X_i \log \frac{\tilde{\pi}}{\hat{\pi}} + (1 - X_i) \log \frac{1 - \tilde{\pi}}{1 - \hat{\pi}} \right\}}.$$


4. Return $p_n^{\text{spaCRT}} = 1 - \Phi(r_n) + \phi(r_n) \left(\frac{1}{\lambda_n} - \frac{1}{r_n} \right).$

Results on Gasperini data analysis


Results on Gasperini data analysis



Results on Gasperini data analysis



Results on Gasperini data analysis



Conclusions




Conclusions

Prospective experimental design: PerturbPlan balances cost and power in Perturb-seq experiments, supporting 11 common design problems.



Conclusions


Prospective experimental design: PerturbPlan balances cost and power in Perturb-seq experiments, supporting 11 common design problems.



Conclusions

Prospective experimental design: PerturbPlan balances cost and power in Perturb-seq experiments, supporting 11 common design problems.

Post-experimental analysis: spaCRT supports computationally efficient differential expression analysis, with strong statistical validity guarantees.




Conclusions

Prospective experimental design: PerturbPlan balances cost and power in Perturb-seq experiments, supporting 11 common design problems.

Post-experimental analysis: spaCRT supports computationally efficient differential expression analysis, with strong statistical validity guarantees.

Both tools are open-source and available for use by the community.



References

- Niu, Z., He, Y., Galante, J., Gschwind, A.R., Engreitz, J.M., Katsevich, E. (2025). PerturbPlan: an analytical framework for designing Perturb-seq experiments. In preparation.
- Niu, Z., Ray Choudhury, J., Katsevich, E. (2025). The conditional saddlepoint approximation for fast and accurate large-scale hypothesis testing. arXiv:2407.08911.
- Gasperini, M. et al. (2019). A genome-wide framework for mapping gene regulation via cellular genetic screens. Cell.
- Replogle, J. et al. (2022). Mapping information-rich genotype-phenotype landscapes with genome-scale Perturb-seq. Cell.
- Barry, T. et al. (2021). SCEPTRE improves calibration and sensitivity in single-cell CRISPR screen analysis. Genome Biology.
- Liu, M., Katsevich, E., Janson, L., and Ramdas, A. (2022). Fast and powerful conditional randomization testing via distillation. Biometrika.
- Shah, R. & Peters, J. (2020). The hardness of conditional independence testing and the GCM test. Annals of Statistics.

Thank you!
Questions?